

KAFKA ADMINISTRATION

Đơn vị: Công ty CP Giáo dục và Công nghệ QNET

QNET JOINT STOCK COMPANY

Address: 14th Floor, VTC Online Tower
18 Tam Trinh Street. Hoang Mai District
Hanoi, Vietnam



Quality Network for Education and Technology


NỘI DUNG

- Introduction, Understanding Topics and Partitions and Brokers
- Kafka Producer, Kafka Consumer
- Kafka Operations and Performance Tuning
- Kafka Cluster Setup & Administration
- Kafka Monitoring and Schema Registry
- Admin Client and Securing Kafka
- Known Issues in Apache Kafka
- Debugging and Troubleshooting Kafka Connect
- Key points and recommendation on Apache Kafka stream processing



KAFKA: Publish/Subscribe Messaging

Single, direct metrics publisher



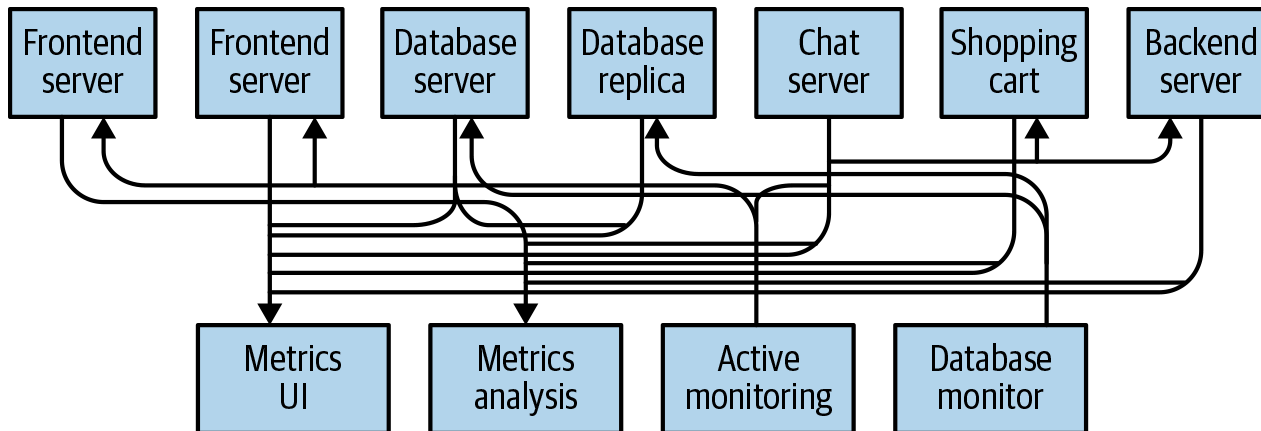
Simple publish/subscribe system with a simple message queue

Or interprocess communication channel:

- Push metrics from frontend server to metrics server to display on its dashboard

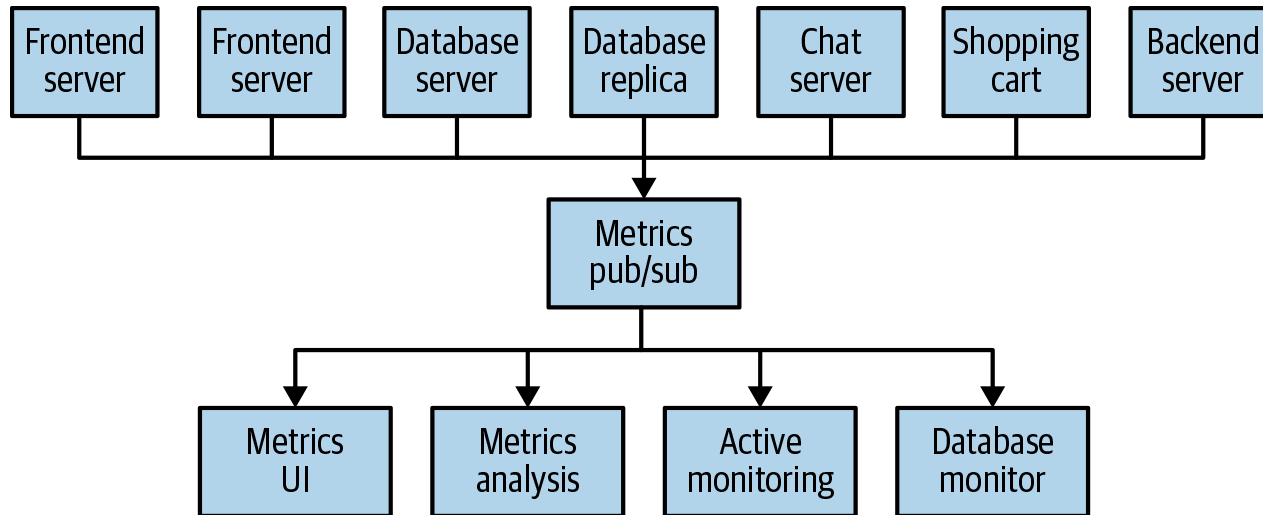
Introduction

Many metrics publishers, using direct connection

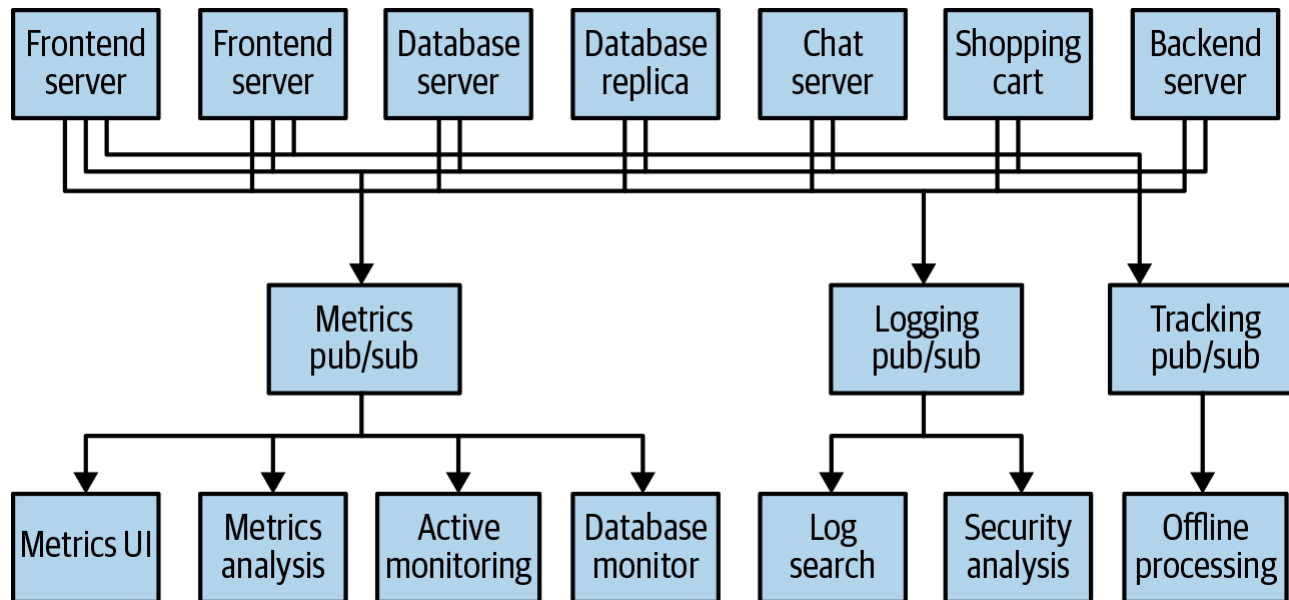


Introduction

Metric publish/subscribe system



Multiple publish/subscribe system



KAFKA

- Apache Kafka is a distributed publish/subscribe messaging system which allows publishing data, which will grow as per your business
- Apache Kafka was developed to solve LinkedIn data pipeline problem
Kafka was created at LinkedIn
 - It was designed to provide a high-performance messaging system that can handle user activity and system metrics in real time
 - Release in 2010 as an Github open source project

Features of Kafka

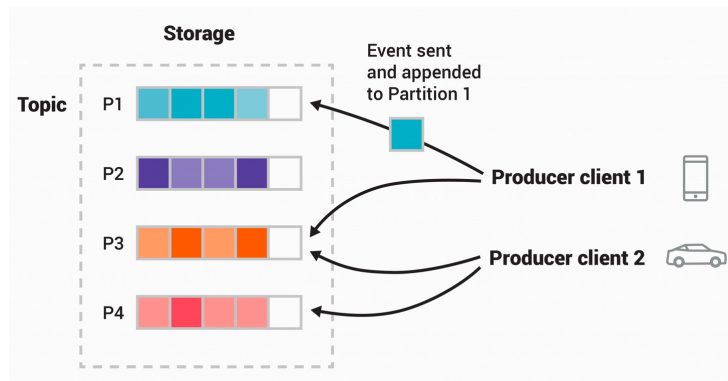
- High throughput
 - Support for millions of messages per second
- Data loss
 - Ensures no data loss
 - Provides compression & Security
- Durability
 - Provides support to persisting messages on disk
- Scalability
 - Highly scalable distributed systems with no downtime
- Stream Processing
 - Kafka can be used along with streaming framework: Spark, Flink, Storm...
- Replication
 - Messages can be replicated across clusters, which supports multiple subscribers

KAFKA USECASE

- Messaging
- Website Activity Tracking
- Metrics
- Log Aggregation
- Stream Processing
- Event Sourcing
- Commit Log

MAIN CONCEPTS AND TERMINOLOGY

Message (Event)

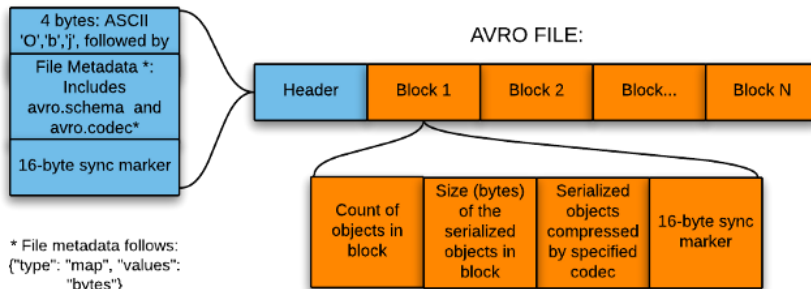


- An event records the fact that ‘something happened’ in the world or in your business
- An event is also called record or message
- When you read and write data to Kafka, you do this in the form of events
- An event has:
 - Key
 - Value
 - Timestamp
 - Optional metadata headers
- Example:
 - Event key: Alice
 - Event Value: “Made a payment of \$200 to Bob”
 - Event timestamp: “Jun, 25,2020 at 2:06 pm”

MAIN CONCEPTS AND TERMINOLOGY

Schema

- JSON
- XML
- Avro
- Protobuf



MAIN CONCEPTS AND TERMINOLOGY

Topics

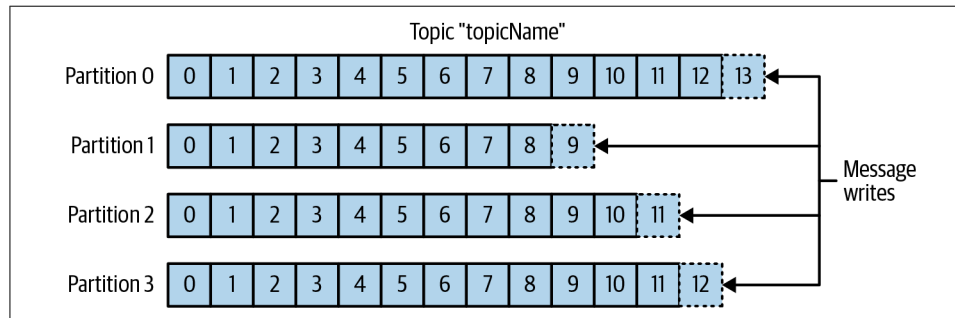


Figure 1-5. Representation of a topic with multiple partitions

- Messages are organized and durably stored in *topics*.
- Topics is similar to a folder in a file system and the events are the files in that folder
- Topics in Kafka are always multi-producers and multi-subscriber
- Messages in a topic can be read as often as needed and are not deleted after consumption unlike traditional messaging systems

MAIN CONCEPTS AND TERMINOLOGY

Partitions

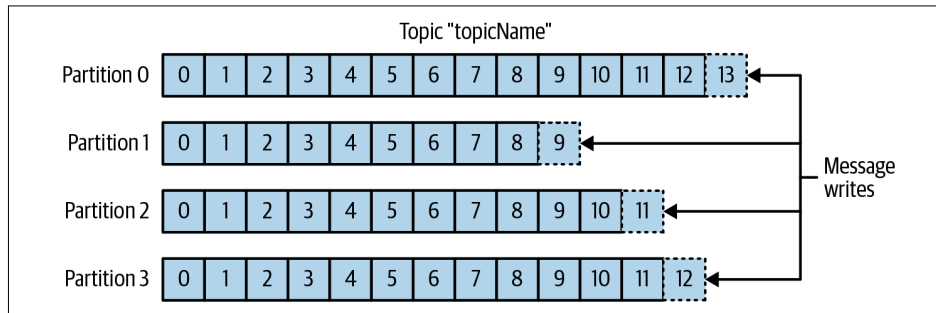
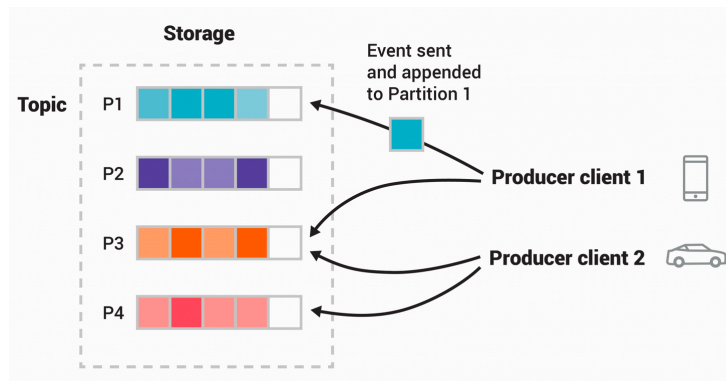


Figure 1-5. Representation of a topic with multiple partitions

- Topics are partitioned, meaning a topic is spread over a number of “bucket” located on different Kafka brokers.
- When a new event is published to a topic, it is actually appended to one of the topic’s partitions.
- Events with the same event key are written to the same partition

MAIN CONCEPTS AND TERMINOLOGY

Producers and Consumers



- Producers are those client applications that publish (write) events to Kafka
- Consumers are those that subscribe to (read and process) these events.

MAIN CONCEPTS AND TERMINOLOGY

Producers and Consumers

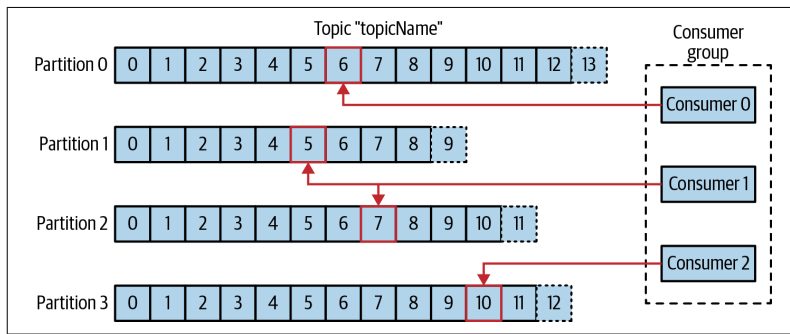
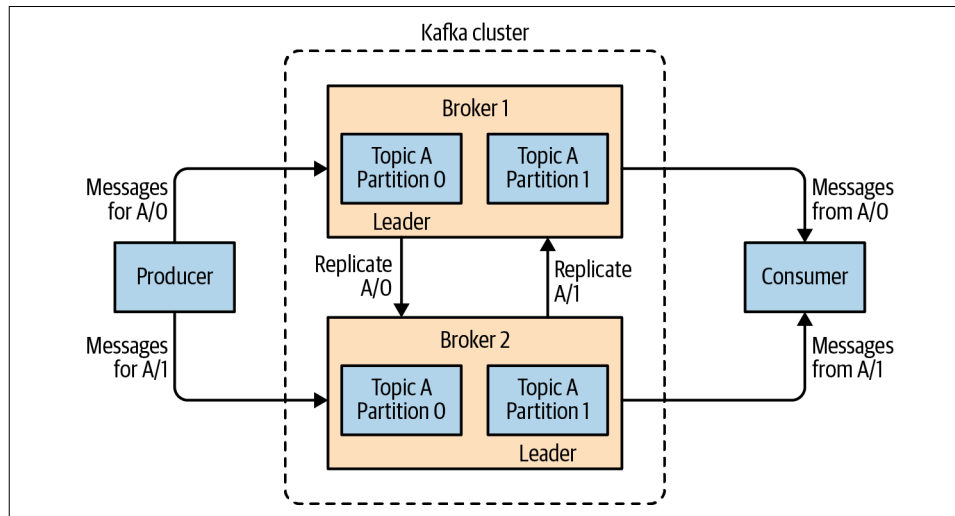


Figure 1-6. A consumer group reading from a topic

- Producers are those client applications that publish (write) events to Kafka
- Consumers are those that subscribe to (read and process) these events.

MAIN CONCEPTS AND TERMINOLOGY

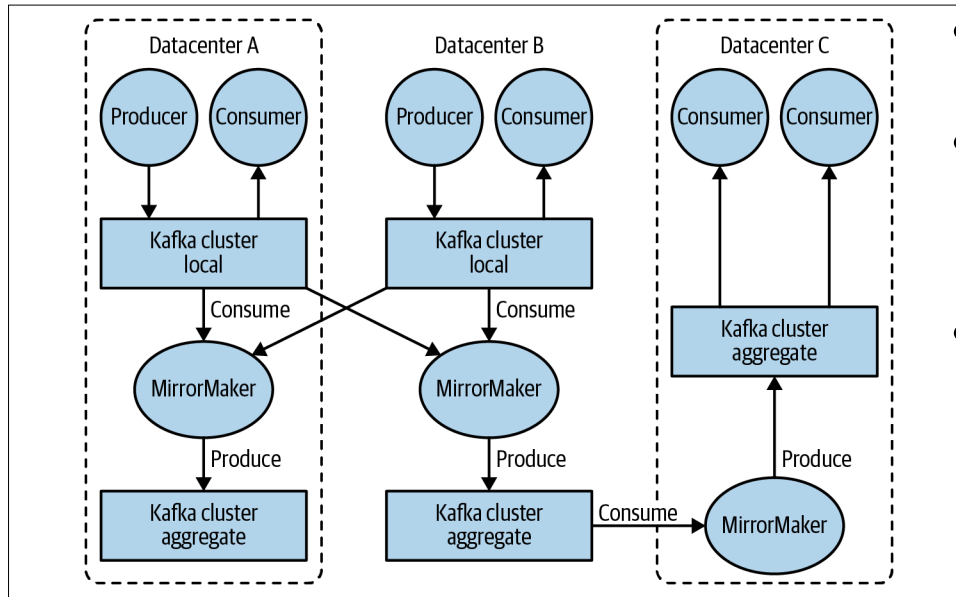
Brokers and Clusters



- A single Kafka server is called a **broker**
- Kafka brokers are designed to operate as part of a **cluster**

MAIN CONCEPTS AND TERMINOLOGY

Multiple Clusters



- Segregation of types of data
- Isolation for security requirements
- Multiple datacenters (DR)

KAFKA API

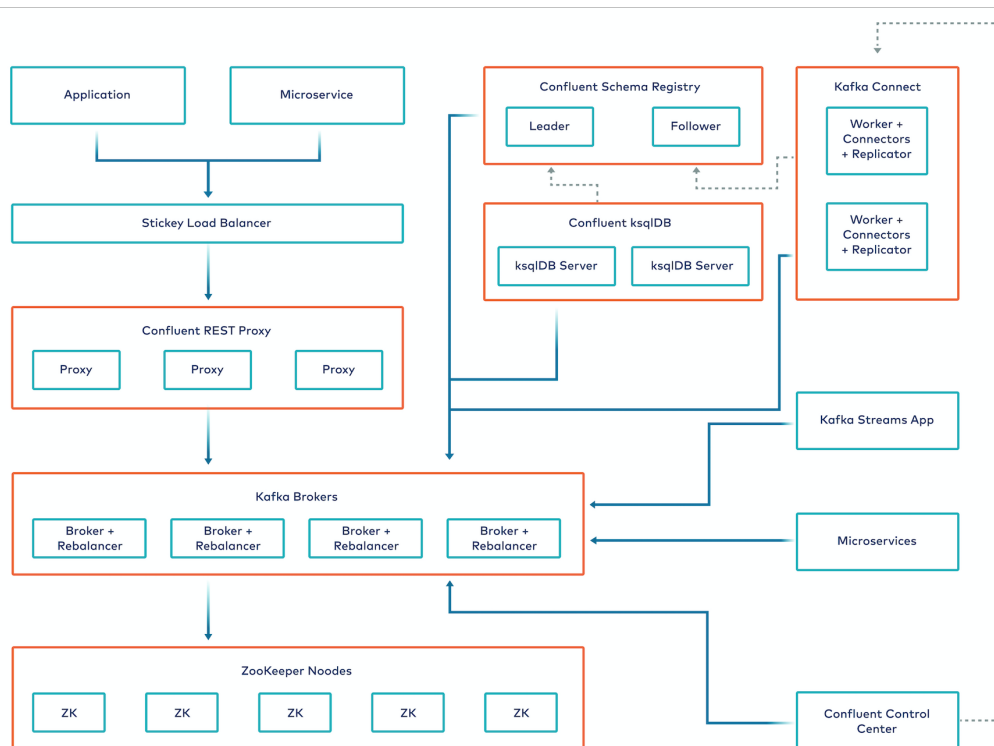
- Admin API: to manage and inspect topics, brokers and other Kafka objects
- Producer API: to publish a stream of events to one or more Kafka topics
- Consumer API: to subscribe one or more topics and to process the stream of events
- Kafka Stream API: to implement stream processing applications and microservices. It provides higher-level functions to process event streams, including transformations, stateful operations like aggregations and joins, windowing....
- Kafka Connect API : to build and run reusable data import/export connectors that consume or producer stream of events

Why Kafka ?

- Multiple Producers
- Multiple Consumers
- Disk-bases retention
- Scalable
- High Performance
- Platform Features

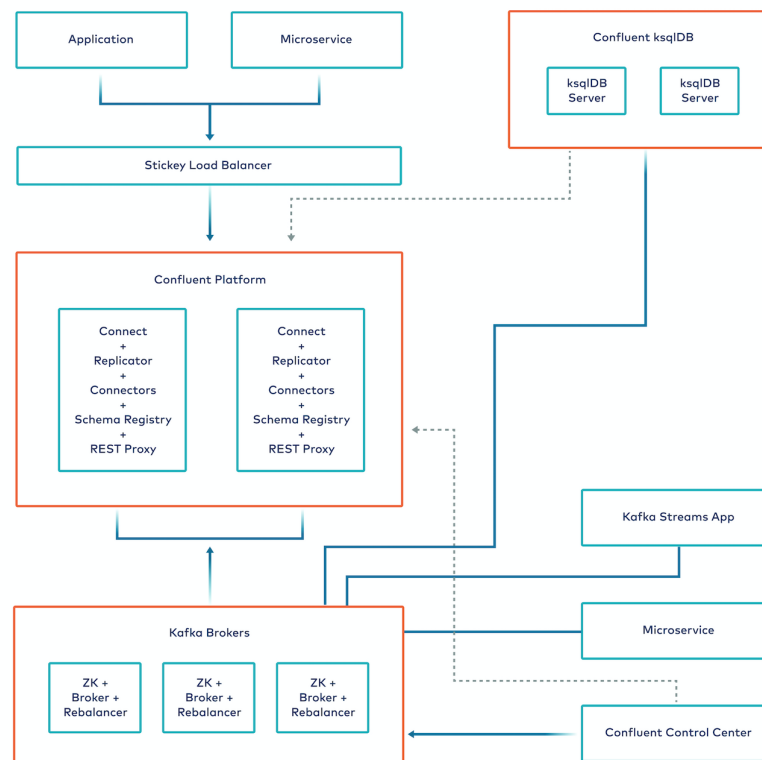
Confluent Platform

Large Cluster Reference Architecture



Confluent Platform

Small Cluster Reference Architecture



CONCLUSION



Quality Network for Education and Technology

XIN CHÂN THÀNH CẢM ƠN!